AI GEN X × DRAGONFRUIT MEDIA

# Implementation Roadmap & Risk Framework

## Batch 6 — Phased Delivery, Risk Mitigation, Governance & Success Metrics

**BATCH 6 OF 6 — FINAL DELIVERABLE**

March 2026 · Confidential

Prepared exclusively for Dragonfruit Media

Evaluation Criteria Alignment: 20% Realistic Roadmap + 10% Engagement Fit = 30% of total score

# Table of Contents


| | | |
|---|---|---|
| 1 | Roadmap Philosophy & Design Principles | 3 |
| 2 | Master Timeline — 6-Month Gantt Overview | 4 |
| 3 | Phase 0 — Foundation & Quick Wins (Days 1–30) | 5 |
| 4 | Phase 1 — Core Agent Infrastructure (Months 2–3) | 8 |
| 5 | Phase 2 — Creative Intelligence Suite (Months 3–4) | 11 |
| 6 | Phase 3 — Full Deployment & Scale (Months 5–6) | 13 |
| 7 | Risk Framework — 14 Risks with Mitigation Strategies | Part 2 |
| 8 | Human-in-the-Loop Governance Model | Part 2 |
| 9 | Training & Adoption Framework | Part 2 |
| 10 | Engagement Model & Working Cadence | Part 2 |
| 11 | Success Metrics & KPI Dashboard | Part 2 |
| 12 | Rollback & Contingency Playbook | Part 2 |
| 13 | Post-Engagement Sustainability Plan | Part 2 |
| 14 | Evaluation Criteria Final Alignment | Part 2 |

# 1. Roadmap Philosophy & Design Principles

> **Core Principle:** Quick wins build trust. Trust enables transformation. Transformation enables 7 clients per pod.

## 1.1 Five Design Principles

| # | Principle | What It Means for DFM | Brief Criteria |
|---|---|---|---|
| 1 | **Quick Wins First** | Visible value within 30 days. Slack templates, Make expansions, and Claude Projects require zero infrastructure — they use tools DFM already pays for. This builds organizational confidence before we introduce agents. | 20% Roadmap |
| 2 | **Pilot → Rollout** | Every agent deploys to **one pod first** (recommended: Mango Pod). We measure, iterate, then expand. No big-bang risk. If an agent underperforms, only one pod is affected. | 20% Roadmap |
| 3 | **Human-in-the-Loop by Default** | All write operations (create task, send message, publish video) require human approval until graduation criteria are met. 7 HITL gates defined in Batch 2, with a graduation path from `ALWAYS_CONFIRM` → `NOTIFY_ONLY` → `AUTONOMOUS`. | 30% Specificity |
| 4 | **Pipeline as Source of Truth** | Every deliverable maps back to DFM's actual pipeline phases (Onboarding → Pre-Pro → Production → Post-Pro → Publish). Aligned with REACT_DEMO_PLAN.md hierarchy: Pipeline → Phase → Role → Opportunity. | 30% Specificity |
| 5 | **Measurable at Every Phase** | Each phase has exit criteria tied to DFM's dream allocations. We don't move to Phase 2 until Phase 1 KPIs are met. No scope creep, no ambiguity. | 20% Roadmap |

## 1.2 Roadmap Alignment to Evaluation Criteria

> **This document directly addresses:**
> • **20% Realistic Roadmap** — Phased rollout, milestones, KPIs, risk mitigation, training plan
> • **10% Engagement Fit** — Async-first, Slack-heavy, biweekly steering, written updates
> • Cross-references **30% Specificity** (pipeline phases, time audit data, tool names) and **25% Architecture** (from Batch 2)
> • Total evaluation criteria coverage across all 6 batches: **100%**

## 1.3 Dependency Chain

| Phase | Depends On | Enables | Risk if Delayed |
|---|---|---|---|
| **Phase 0** (Quick Wins) | Nothing — starts Day 1 | Organizational trust, API access provisioning | Low — independent deliverables |
| **Phase 1** (Core Agents) | API keys (ClickUp, Slack), Phase 0 trust | PR-4, PR-3, ED-2, ED-1 agents | Medium — blocks all agent work |
| **Phase 2** (Creative Suite) | Phase 1 agent framework, Notion API | CD-1 through CD-9, PR-1, PR-2 | Medium — CD tools delayed |
| **Phase 3** (Scale) | Phase 1+2 pilot success, all API access | Multi-pod rollout, 7 clients/pod | High — delays EOY target |

# 2. Master Timeline — 6-Month Gantt Overview

## 2.1 High-Level Gantt

| Workstream | M1 | M2 | M3 | M4 | M5 | M6 | Key Milestone |
|---|---|---|---|---|---|---|---|
| **Phase 0: Quick Wins** | | | | | | | Slack templates live, Make expanded, Claude Projects for 3 clients |
| **Infrastructure** | | | | | | | OpenAI Agents SDK deployed, MCP servers running |
| **Phase 1: Core Agents** | | | | | | | PR-4, PR-3, ED-2, ED-1 live on pilot pod |
| **Phase 2: Creative Suite** | | | | | | | CD-1 through CD-5, PR-1, PR-2 live |
| **Phase 3: Scale** | | | | | | | All 16 tools deployed, multi-pod, 7 clients/pod |
| **Training** | | | | | | | Continuous — role-specific, phase-aligned |
| **React Demo** | | | | | | | Interactive demo aligned with REACT_DEMO_PLAN.md |

## 2.2 Milestone Calendar

| Day/Month | Milestone | Deliverable | Exit Criteria |
|---|---|---|---|
| **Day 1** | Engagement kickoff | Shared Slack channel, repo access, API key requests sent | Channel active |
| **Day 7** | Slack templates deployed | PR-3: 12 templates across 4 categories in Slack | Templates in use |
| **Day 14** | Make scenarios expanded | 5 new Make scenarios (from Batch 4 inventory) | Scenarios running |
| **Day 21** | Claude Projects live | CD-4: Client Knowledge Bases for 3 pilot clients | CDs using daily |
| **Day 30** | Phase 0 complete ✅ | | KPIs met (below) |

| | | | |
|---|---|---|---|
| | | Quick wins report, Phase 1 readiness assessment | |
| **Month 2** | First agent live | PR-4 Claude-to-ClickUp on pilot pod | 50+ tasks created via agent |
| **Month 2** | PC agent live | ED-2 Capacity-Based Assignment on pilot pod | Assignments match manual |
| **Month 3** | QC + Sentiment live | ED-1 Video QC Checker, ED-5 Sentiment Analyzer | PPM utilization ≤105% |
| **Month 3** | Phase 1 complete ✅ | 4 core agents on pilot pod, HITL gates active | All exit criteria met |
| **Month 4** | Creative suite pilot | CD-1 Retention, CD-2 Competitor, CD-3 Roaster, CD-5 Voice | CDs using 3+ tools weekly |
| **Month 4** | Producer analytics | PR-1 Timeline Calculator, PR-2 Client Sentiment | Producers report time savings |
| **Month 5** | Multi-pod rollout begins | Phase 1+2 agents deployed to Pod 2 and Pod 3 | Rollout playbook followed |
| **Month 6** | Full deployment ✅ | All 16 wish-list tools live, all pods | 7 clients/pod (≥1 pod) |
| **EOY 2026** | Target achieved | Sustained 7 clients/pod across multiple pods | Dream allocations ≥60% |

# 3. Phase 0 — Foundation & Quick Wins (Days 1–30)

**PHASE 0** · Days 1–30 · Zero Infrastructure Required · Immediate Value

> **Why Phase 0 matters:** The brief asks for "quick wins (first 30 days)." These deliverables use tools DFM already pays for ($12,296/mo) — no new infrastructure, no API keys, no agent framework. Pure value extraction from existing spend.

## 3.1 Week 1: Slack Template Library (PR-3)

| Attribute | Detail |
|---|---|
| **Agent ID** | PR-3 — Slack Communication Templates |
| **Roles Served** | Producers (primary), PCs, PPMs |
| **Current State** | Producers spend 17% of time on Slack (time audit). Messages are ad-hoc, inconsistent across pods. |
| **Deliverable** | 12 Slack Workflow Builder templates across 4 categories |
| **Dream Impact** | Slack 17% → 10% (Producers). 7% freed = 3.5 hrs/week/producer |

### Template Categories

| Category | Templates | Trigger |
|---|---|---|
| **Client Updates** | • Batch kickoff notification<br>• Weekly status summary<br>• Milestone completion alert | Slash command: `/update [client] [type]` |
| **Internal Handoffs** | • Script-to-edit handoff<br>• Edit-to-review handoff<br>• Review-to-publish handoff | ClickUp status change → Make → Slack |
| **Feedback Routing** | • Client feedback received<br>• Revision request<br>• Approval confirmation | Frame.io webhook → Make → Slack |
| **Standup/Reporting** | • Daily standup prompt<br>• Weekly capacity report<br>• Monthly client summary | Scheduled: daily 9am, weekly Monday, monthly 1st |

### Implementation Steps

1. **Day 1–2:** Audit existing Slack messages across 3 pods. Identify top 12 recurring message patterns.

2. **Day 3–4:** Build templates in Slack Workflow Builder. Configure slash commands and triggers.
3. **Day 5:** Deploy to pilot pod. Train Producers and PCs (15-min walkthrough).
4. **Day 6–7:** Collect feedback, iterate. Deploy to remaining pods.

## 3.2 Week 2: Make/Zapier Expansion

| Attribute | Detail |
|---|---|
| **Current State** | 7 existing Make scenarios (inventoried in Batch 4). $693/mo spend on Make + Zapier. |
| **Deliverable** | 5 new Make scenarios extending existing infrastructure |
| **Dream Impact** | ClickUp Administration: Producers 38% → 28% (10% reduction from automation alone) |

### New Make Scenarios

| # | Scenario | Trigger | Actions | Time Saved |
|---|---|---|---|---|
| 1 | **Auto-Status Sync** | ClickUp status change | Update linked tasks, notify assignee in Slack, log in Notion | 2 hrs/wk |
| 2 | **Footage Ingest Alert** | New file in Dropbox client folder | Create ClickUp task, notify PC in Slack, update Notion tracker | 1.5 hrs/wk |
| 3 | **Review Deadline Nudge** | ClickUp due date - 24hrs | Send Slack reminder to reviewer, escalate if overdue by 48hrs | 1 hr/wk |
| 4 | **Client Feedback Router** | Frame.io comment webhook | Parse comment, create ClickUp subtask, notify editor in Slack | 3 hrs/wk |
| 5 | **Weekly Capacity Report** | Scheduled: Monday 8am | Pull ClickUp task counts per editor, calculate load, post to #capacity channel | 2 hrs/wk |

> **Total Phase 0 Make savings:** 9.5 hrs/week across all pods = **494 hrs/year** = $14,820 capacity value (at $30/hr blended rate)

## 3.3 Week 3: Claude Projects for Creative Directors (CD-4)

| Attribute | Detail |
|---|---|
| **Agent ID** | CD-4 — Client Knowledge Bases |
| **Current State** | CDs: 100% use ChatGPT (survey). Knowledge scattered across Notion, Slack, Google Docs. CDs spend time re-explaining client context to AI tools. |

| | |
|---|---|
| **Deliverable** | Claude Projects configured for 3 pilot clients with uploaded context |
| **Dream Impact** | CD preproduction time: 29% → 35% (immediate 6% gain from faster ideation) |

### Per-Client Claude Project Structure

| Knowledge Category | Source | Upload Format |
|---|---|---|
| Channel strategy & positioning | Notion client page | Markdown export |
| Top 20 performing videos (titles, hooks, retention) | YouTube Analytics export | CSV |
| Brand voice guidelines | Client onboarding docs | PDF |
| Competitor landscape | CD research notes | Markdown |
| Past script examples (top 5) | Google Docs | Text files |
| Client feedback patterns | Slack history export | Text summary |
| Thumbnail style preferences | Notion mood boards | Image descriptions + links |

### Implementation Steps

1. **Day 15–16:** Select 3 pilot clients (highest video volume). Gather knowledge artifacts.
2. **Day 17–18:** Configure Claude Projects. Upload knowledge bases. Write system prompts per client.
3. **Day 19:** Train CDs (30-min session). Demonstrate ideation, scripting, and packaging workflows.
4. **Day 20–21:** CDs use for one batch cycle. Collect feedback.

## 3.4 Week 4: API Access & Phase 1 Readiness

| # | Task | Owner | Status Gate |
|---|---|---|---|
| 1 | ClickUp API key + Team ID provisioned | DFM IT / Vhinia | REQUIRED |
| 2 | Slack Bot Token created (scopes: chat:write, channels:read, commands) | DFM IT | REQUIRED |
| 3 | Notion API integration created | DFM IT | PHASE 2 |
| 4 | Frame.io API token provisioned | DFM IT / Mahmoud | PHASE 2 |
| 5 | YouTube Data API key | DFM IT | PHASE 2 |
| 6 | OpenAI API key (AI Gen X provides) | AI Gen X | REQUIRED |
| 7 | Pilot pod selected (recommend: Mango Pod) | Brendan / Yazan | REQUIRED |
| 8 | DFM stakeholder designated | Leadership | REQUIRED |

| | | | |
|---|---|---|---|
| 9 | Shared GitHub repo created | AI Gen X | DAY 1 |
| 10 | Dedicated Slack channel (#ai-automation) | AI Gen X + DFM | DAY 1 |

## 3.5 Phase 0 Exit Criteria

> **Phase 0 is complete when ALL of the following are true:**

| # | Criterion | Measurement | Target |
|---|---|---|---|
| 1 | Slack templates deployed and in use | Template usage count in Slack analytics | ≥50 uses/week |
| 2 | Make scenarios running without errors | Make execution logs | ≥95% success rate |
| 3 | Claude Projects active for 3 clients | CD self-report + usage logs | Daily use by ≥2 CDs |
| 4 | API keys provisioned for Phase 1 | Checklist above | Items 1,2,6,7,8 ✅ |
| 5 | Pilot pod selected and briefed | Kickoff meeting held | Pod lead confirmed |
| 6 | Coordination overhead reduction visible | Producer self-report | ≥5% reduction felt |

# 4. Phase 1 — Core Agent Infrastructure (Months 2–3)

**PHASE 1** · Months 2–3 · OpenAI Agents SDK · Pilot Pod · 4 Core Agents

> **Phase 1 Goal:** Deploy the 4 highest-ROI agents on the pilot pod. Establish the agent framework, MCP servers, HITL gates, and observability infrastructure that all subsequent agents will use.

## 4.1 Infrastructure Sprint (Month 2, Week 1–2)

| # | Deliverable | Technical Detail | Days |
|---|---|---|---|
| 1 | **Agent Framework** | OpenAI Agents SDK setup: `Runner`, `Agent`, `FunctionTool` classes. Environment config. Tracing enabled. | 3 |
| 2 | **ClickUp MCP Server** | MCP server exposing: `list_lists`, `create_task`, `update_task_status`, `get_task`, `add_comment`, `search_tasks`. OAuth2 auth. | 3 |
| 3 | **Slack MCP Server** | MCP server exposing: `send_message`, `list_channels`, `get_channel_history`, `add_reaction`. Bot token auth. | 2 |
| 4 | **HITL Gateway** | Approval middleware: intercepts write operations, sends Slack approval request, waits for ✅/❌, logs decision. Configurable per-agent, per-tool. | 2 |
| 5 | **Demo Mode** | Mock data layer (aligned with REACT_DEMO_PLAN.md): `pipeline_stages.json`, `wishlist_tools.json`, `time_allocations.json`, `ai_survey_summary.json`. All connectors return mock data when `DEMO_MODE=true`. | 2 |
| 6 | **Observability** | OpenAI SDK tracing + custom spans. PostgreSQL trace store. Dashboard: agent runs, tool calls, latency, errors, HITL approval rates. | 2 |

> **Infrastructure Sprint Total:** 14 engineering days. Parallel work possible — framework + MCP servers can be built simultaneously by 2 engineers.

## 4.2 Agent Deployment Schedule

| Agent | Name | Role | Deploy | Tools Bound | Annual ROI |
|---|---|---|---|---|---|
| **PR-4** | Claude-to-ClickUp | Producer | Month 2, Wk 3 | ClickUp MCP ( `create_task` , `list_lists` ), Slack MCP ( `send_message` ) | **$72,280** |
| **PR-3** | Slack Templates (Agent) | Producer | Month 2, Wk 3 | Slack MCP ( `send_message` with template selection) | **$31,200** |
| **ED-2** | Capacity Assignment | PC/PPM | Month 2, Wk 4 | ClickUp MCP ( `search_tasks` , `update_task` ), Slack MCP ( `send_message` ) | **$31,200** |
| **ED-1** | Video QC Checker | PPM | Month 3, Wk 1 | Frame.io API ( `get_asset` , `add_comment` ), ClickUp MCP ( `update_task` ) | **$72,280** |
| **ED-5** | Sentiment Analyzer | PPM | Month 3, Wk 2 | Frame.io API ( `list_comments` ), Slack MCP ( `send_message` ) | **$31,200** |

## 4.3 Per-Agent Deployment Playbook

Each agent follows the same 5-step deployment process:

| Step | Activity | Duration | Detail |
|---|---|---|---|
| 1 | **Build** | 3–5 days | Agent code, tool bindings, instructions, guardrails. Unit tests with mock data. |
| 2 | **Demo Mode Test** | 1 day | Run against mock data. Verify ReAct flow, tool calls, output format. Record agent trajectory GIF for React demo. |
| 3 | **Pilot Deploy** | 1 day | Deploy to pilot pod with `HITL_MODE=ALWAYS_CONFIRM` . All write operations require Slack approval. |
| 4 | **Pilot Run** | 5–10 days | Monitor: approval rate, error rate, user satisfaction. Iterate on agent instructions based on feedback. |
| 5 | **Graduate** | 1 day | If approval rate ≥95% over 50+ runs: graduate to `NOTIFY_ONLY` . If <90%: iterate and re-pilot. |

## 4.4 PR-4 Detailed Deployment (Flagship Agent)

PR-4 (Claude-to-ClickUp) is the flagship agent — highest ROI ($72,280/yr), most visible, and sets the pattern for all subsequent agents.

### Week-by-Week Plan

| Timeline | Activity | Exit Criteria |
|---|---|---|
| **M2 W3 D1-3** | Build PR-4 agent: instructions (from Batch 3 playbook), bind ClickUp MCP + Slack MCP tools, implement ReAct flow, add guardrails (max 3 tasks/run, list validation) | Agent creates correct mock tasks from natural language |
| **M2 W3 D4** | Demo mode test: run 20 test scenarios (from Batch 3 test cases). Record agent trajectory for React demo GIF. | 18/20 scenarios pass. GIF recorded. |
| **M2 W3 D5** | Deploy to pilot pod. Train pilot Producer (15-min walkthrough). HITL = ALWAYS_CONFIRM. | Producer can invoke agent from Slack or dashboard |
| **M2 W4** | Pilot run: Producer uses PR-4 for all new task creation. AI Gen X monitors traces daily. Iterate instructions based on edge cases. | 50+ tasks created. Approval rate tracked. |
| **M3 W1** | Graduation assessment: If approval rate ≥95%, graduate to NOTIFY_ONLY. If not, iterate one more week. | HITL mode updated. Producer confirms satisfaction. |

## 4.5 Cross-Agent Handoff Testing (Month 3)

Once PR-4 and ED-2 are both live, we test the first cross-agent handoff:

> **Handoff Test:** Producer creates batch tasks via PR-4 → PR-4 triggers ED-2 → ED-2 assigns editors based on capacity → ED-2 notifies editors via Slack.
>
> **Success criteria:** End-to-end flow completes in <30 seconds. All HITL gates fire correctly. No orphaned tasks.

## 4.6 Phase 1 Exit Criteria

| # | Criterion | Measurement | Target |
|---|---|---|---|
| 1 | PR-4 live on pilot pod | Agent trace logs | ≥50 tasks created |
| 2 | ED-2 live on pilot pod | Assignment accuracy vs. manual | ≥90% match |
| 3 | ED-1 QC Checker operational | Videos checked automatically | ≥20 videos checked |
| 4 | HITL gates functioning | Approval/rejection logs | 100% of writes gated |
| 5 | Producer ClickUp time reduced | Time audit comparison | 38% → ≤30% |
| 6 | PPM utilization improved | Time audit comparison | 121% → ≤110% |
| 7 | Cross-agent handoff working | PR-4 → ED-2 end-to-end test | ≥10 successful handoffs |

| | | | |
|---|---|---|---|
| 8 | Observability dashboard live | Dashboard accessible | All metrics visible |

# 5. Phase 2 — Creative Intelligence Suite (Months 3–4)

**PHASE 2** · Months 3–4 · CD & Producer Analytics · Notion Integration · 8 New Agents

> **Phase 2 Goal:** Deploy the Creative Director toolkit (CD-1 through CD-5, CD-9) and Producer analytics (PR-1, PR-2). This phase transforms CDs from "100% ChatGPT" (survey) to a purpose-built creative intelligence suite.

## 5.1 Agent Deployment Schedule

| Agent | Name | Role | Deploy | Tools Bound | ROI/yr |
|---|---|---|---|---|---|
| **CD-1** | Retention Scoring | CD | M3 W3 | YouTube API ( `get_analytics` ), Notion MCP ( `create_page` ) | $31,200 |
| **CD-2** | Competitor Scraper | CD | M3 W4 | YouTube API ( `search` , `get_video_stats` ), Notion MCP ( `create_page` ) | $31,200 |
| **CD-3** | Packaging Roaster | CD | M4 W1 | Slack MCP ( `/roast` command), YouTube API | $15,600 |
| **CD-5** | Voice Matching | CD | M4 W1 | Notion MCP ( `get_page` ), Claude API (style analysis) | $15,600 |
| **CD-9** | Outlier Detection | CD | M4 W2 | YouTube API ( `get_channel_videos` ), Slack MCP | $15,600 |
| **PR-1** | Timeline Calculator | Producer | M4 W2 | ClickUp MCP ( `get_task` , `search_tasks` ), historical data | $15,600 |
| **PR-2** | Client Sentiment | Producer | M4 W3 | Slack MCP ( `get_channel_history` ), sentiment model | $15,600 |
| **CD-4+** | Knowledge Base (Agent) | CD | M4 W3 | Notion MCP (full CRUD), vector search, Claude API | $15,600 |

## 5.2 Notion MCP Server Build

| Attribute | Detail |
|---|---|
| **Timeline** | Month 3, Week 3 (parallel with CD-1 build) |
| **Tools Exposed** | |

| | |
|---|---|
| | `get_page`, `create_page`, `update_page`, `search`, `query_database`, `append_block` |
| **Auth** | Notion Internal Integration token. Scoped to client knowledge base databases only. |
| **Used By** | CD-1, CD-2, CD-4, CD-5, CD-9 (5 agents) |
| **Build Effort** | 2 engineering days (official Notion MCP available) |

## 5.3 CD Workflow Transformation

Phase 2 transforms the Creative Director workflow from manual research to AI-augmented creative intelligence:

| Current (Survey) | → | Phase 2 State | Enabling Agents |
|---|---|---|---|
| Manual YouTube research (2+ hrs/client/week) | → | Automated competitor monitoring | CD-2 Competitor Scraper, CD-9 Outlier Detection |
| Gut-feel retention predictions | → | Data-driven retention scoring | CD-1 Retention Scoring |
| Ad-hoc ChatGPT for ideation | → | Client-context-aware ideation | CD-4 Knowledge Base (Agent), CD-5 Voice Matching |
| Subjective packaging review | → | Quantitative packaging analysis | CD-3 Packaging Roaster |

> **Phase 2 Dream Impact:** CD preproduction time 29% → 50%. This is the single largest creative quality improvement in the entire engagement. CDs spend less time on ClickUp (15% → 0%), V1 Review (15% → 5%), and Post-Pub Monitoring (12% → 0%), freeing 21% for ideation, scripting, and strategic creative work.

## 5.4 Phase 2 Exit Criteria

| # | Criterion | Measurement | Target |
|---|---|---|---|
| 1 | CD-1 through CD-5 live on pilot pod | Agent trace logs | ≥5 runs/agent/week |
| 2 | PR-1 and PR-2 live on pilot pod | Agent trace logs | Producers using weekly |
| 3 | Notion MCP server operational | Uptime monitoring | ≥99% uptime |

| | | | |
|---|---|---|---|
| 4 | CD preproduction time increased | Time audit comparison | 29% → ≥40% |
| 5 | CDs report creative quality improvement | Survey / self-report | ≥80% positive |
| 6 | All Phase 1 agents still healthy | Error rate, usage | <2% error rate |
| 7 | Cross-agent handoffs expanded | Handoff logs | ≥3 handoff patterns active |

# 6. Phase 3 — Full Deployment & Scale (Months 5–6)

**PHASE 3** · Months 5–6 · Multi-Pod Rollout · Remaining Agents · 7 Clients/Pod

**Phase 3 Goal:** Roll out all agents to all pods. Deploy remaining wish-list tools (CD-6 through CD-8, ED-3, ED-4). Achieve 7 clients per pod on at least one pod. Transition to steady-state operations.

## 6.1 Multi-Pod Rollout Strategy

| Week | Pod | Action | Risk Mitigation |
|---|---|---|---|
| **M5 W1** | Pod 2 | Deploy Phase 1 agents (PR-4, PR-3, ED-2, ED-1, ED-5). HITL = ALWAYS_CONFIRM. | Pilot pod Producer mentors Pod 2 Producer. AI Gen X monitors traces. |
| **M5 W2** | Pod 3 | Deploy Phase 1 agents. Parallel with Pod 2 Phase 2 deployment. | Staggered by 1 week to avoid support overload. |
| **M5 W3** | Pod 2 | Deploy Phase 2 agents (CD-1 through CD-5, PR-1, PR-2). | CD from pilot pod runs 30-min training for Pod 2 CD. |
| **M5 W4** | Pod 3 | Deploy Phase 2 agents. | Rollout playbook now battle-tested across 2 pods. |
| **M6 W1** | All Pods | Deploy remaining agents: CD-6, CD-7, CD-8, ED-3, ED-4. | Lower-priority agents. Failures don't block 7-client target. |
| **M6 W2** | Pilot Pod | Add 5th, 6th, 7th client to pilot pod. | Gradual: 1 client/week. Monitor capacity metrics daily. |
| **M6 W3-4** | All Pods | Scale remaining pods to 5-6 clients | |